

## EDA on census income dataset

```
[2]: import pandas as pd
import seaborn as sns
from matplotlib import pyplot as plt
```

```
[3]: df = pd.read_csv(r'C:\Users\Fatman\Documents\adult.data.csv')
```

```
[4]: df
```

```
[4]:
```

|       | age | workclass        | fnlwgt | education  | education-num | \ |
|-------|-----|------------------|--------|------------|---------------|---|
| 0     | 39  | State-gov        | 77516  | Bachelors  | 13            |   |
| 1     | 50  | Self-emp-not-inc | 83311  | Bachelors  | 13            |   |
| 2     | 38  | Private          | 215646 | HS-grad    | 9             |   |
| 3     | 53  | Private          | 234721 | 11th       | 7             |   |
| 4     | 28  | Private          | 338409 | Bachelors  | 13            |   |
| ...   | ... | ...              | ...    | ...        | ...           |   |
| 32556 | 27  | Private          | 257302 | Assoc-acdm | 12            |   |
| 32557 | 40  | Private          | 154374 | HS-grad    | 9             |   |
| 32558 | 58  | Private          | 151910 | HS-grad    | 9             |   |
| 32559 | 22  | Private          | 201490 | HS-grad    | 9             |   |
| 32560 | 52  | Self-emp-inc     | 287927 | HS-grad    | 9             |   |

|       | marital-status     | occupation        | relationship  | race  | sex    | \ |
|-------|--------------------|-------------------|---------------|-------|--------|---|
| 0     | Never-married      | Adm-clerical      | Not-in-family | White | Male   |   |
| 1     | Married-civ-spouse | Exec-managerial   | Husband       | White | Male   |   |
| 2     | Divorced           | Handlers-cleaners | Not-in-family | White | Male   |   |
| 3     | Married-civ-spouse | Handlers-cleaners | Husband       | Black | Male   |   |
| 4     | Married-civ-spouse | Prof-specialty    | Wife          | Black | Female |   |
| ...   | ...                | ...               | ...           | ...   | ...    |   |
| 32556 | Married-civ-spouse | Tech-support      | Wife          | White | Female |   |
| 32557 | Married-civ-spouse | Machine-op-inspct | Husband       | White | Male   |   |
| 32558 | Widowed            | Adm-clerical      | Unmarried     | White | Female |   |
| 32559 | Never-married      | Adm-clerical      | Own-child     | White | Male   |   |
| 32560 | Married-civ-spouse | Exec-managerial   | Wife          | White | Female |   |

|   | capital-gain | capital-loss | hours-per-week | native-country | salary |
|---|--------------|--------------|----------------|----------------|--------|
| 0 | 2174         | 0            | 40             | United-States  | <=50K  |
| 1 | 0            | 0            | 13             | United-States  | <=50K  |
| 2 | 0            | 0            | 40             | United-States  | <=50K  |

|       |       |     |     |               |       |
|-------|-------|-----|-----|---------------|-------|
| 3     | 0     | 0   | 40  | United-States | <=50K |
| 4     | 0     | 0   | 40  | Cuba          | <=50K |
| ...   | ...   | ... | ... | ...           | ...   |
| 32556 | 0     | 0   | 38  | United-States | <=50K |
| 32557 | 0     | 0   | 40  | United-States | >50K  |
| 32558 | 0     | 0   | 40  | United-States | <=50K |
| 32559 | 0     | 0   | 20  | United-States | <=50K |
| 32560 | 15024 | 0   | 40  | United-States | >50K  |

[32561 rows x 15 columns]

```
[5]: df.shape
```

```
[5]: (32561, 15)
```

```
[6]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32561 entries, 0 to 32560
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                    32561 non-null  int64
1   workclass              32561 non-null  object
2   fnlwgt                 32561 non-null  int64
3   education              32561 non-null  object
4   education-num          32561 non-null  int64
5   marital-status         32561 non-null  object
6   occupation              32561 non-null  object
7   relationship            32561 non-null  object
8   race                   32561 non-null  object
9   sex                    32561 non-null  object
10  capital-gain            32561 non-null  int64
11  capital-loss            32561 non-null  int64
12  hours-per-week          32561 non-null  int64
13  native-country          32561 non-null  object
14  salary                  32561 non-null  object
dtypes: int64(6), object(9)
memory usage: 3.7+ MB
```

```
[7]: df.describe()
```

```
[7]:
```

|       | age          | fnlwgt       | education-num | capital-gain | capital-loss \ |
|-------|--------------|--------------|---------------|--------------|----------------|
| count | 32561.000000 | 3.256100e+04 | 32561.000000  | 32561.000000 | 32561.000000   |
| mean  | 38.581647    | 1.897784e+05 | 10.080679     | 1077.648844  | 87.303830      |
| std   | 13.640433    | 1.055500e+05 | 2.572720      | 7385.292085  | 402.960219     |
| min   | 17.000000    | 1.228500e+04 | 1.000000      | 0.000000     | 0.000000       |
| 25%   | 28.000000    | 1.178270e+05 | 9.000000      | 0.000000     | 0.000000       |

|     |           |              |           |              |             |
|-----|-----------|--------------|-----------|--------------|-------------|
| 50% | 37.000000 | 1.783560e+05 | 10.000000 | 0.000000     | 0.000000    |
| 75% | 48.000000 | 2.370510e+05 | 12.000000 | 0.000000     | 0.000000    |
| max | 90.000000 | 1.484705e+06 | 16.000000 | 99999.000000 | 4356.000000 |

|       | hours-per-week |
|-------|----------------|
| count | 32561.000000   |
| mean  | 40.437456      |
| std   | 12.347429      |
| min   | 1.000000       |
| 25%   | 40.000000      |
| 50%   | 40.000000      |
| 75%   | 45.000000      |
| max   | 99.000000      |

1. How many men and women (sex feature) are represented in this dataset?

```
[8]: df['sex'].value_counts()
```

```
[8]: sex
Male      21790
Female    10771
Name: count, dtype: int64
```

2. What is the average age (age feature) of women?

```
[9]: filt = (df['sex'] == 'Female')
avg_a_f = df.loc[filt, 'age'].mean()
print(f"The average age of women is {round(avg_a_f)}.")
```

The average age of women is 37.

3. What is the percentage of German citizens (native-country feature)?

```
[10]: filt = df['native-country'] == 'Germany'
p_of_g = df.loc[filt, 'native-country'].count() / df.shape[0]
print(f"The percentage of German citizens is {p_of_g}.")
```

The percentage of German citizens is 0.004207487485028101.

4-5. What are mean value and standard deviation of the age of those who receive more than 50K per year (salary feature) and those who receive less than 50K per year?

```
[11]: # Approach 1:

s1 = df.loc[df['salary'] == '>50K', 'age']
s2 = df.loc[df['salary'] == '<=50K', 'age']
print(f"The mean value and standard deviation of the age of those who receive_
↳ more than 50K per year (rich) is {round(s1.mean())} and {round(s1.std())}.")
print(f"The mean value and standard deviation of the age of those who receive_
↳ less than 50K per year (poor) is {round(s2.mean())} and {round(s2.std())}.")
```

The mean value and standard deviation of the age of those who receive more than 50K per year(rich) is 44 and 11.

The mean value and standard deviation of the age of those who receive less than 50K per year(poor) is 37 and 14.

[12]: *# Approach 2:*

```
sal_group = df.groupby('salary')[ 'age'].agg(['mean', 'std'])
sal_group
```

```
[12]:
```

|        | mean      | std       |
|--------|-----------|-----------|
| salary |           |           |
| <=50K  | 36.783738 | 14.020088 |
| >50K   | 44.249841 | 10.519028 |

6. Is it true that people who earn more than 50K have at least high school education? (education – Bachelors, Prof-school, Assoc-acdm, Assoc-voc, Masters or Doctorate feature)

```
[13]: filt = (df['salary'] == '>50K')
print(df.loc[filt, 'education'].unique())
print("\nNo, it's not true that people who earn more than 50K have at least_
      ↪high school education.")
```

```
['HS-grad' 'Masters' 'Bachelors' 'Some-college' 'Assoc-voc' 'Doctorate'
 'Prof-school' 'Assoc-acdm' '7th-8th' '12th' '10th' '11th' '9th' '5th-6th'
 '1st-4th']
```

No, it's not true that people who earn more than 50K have at least high school education.

7. Display age statistics for each race (race feature) and each gender (sex feature). Use groupby() and describe(). Find the maximum age of men of Amer-Indian-Eskimo race.

```
[14]: grp_r_s = df.groupby(['race', 'sex'])
      grp_r_s['age'].describe()
```

```
[14]:
```

|                    |        | count   | mean      | std       | min  | 25%  | 50%  | \ |
|--------------------|--------|---------|-----------|-----------|------|------|------|---|
| race               | sex    |         |           |           |      |      |      |   |
| Amer-Indian-Eskimo | Female | 119.0   | 37.117647 | 13.114991 | 17.0 | 27.0 | 36.0 |   |
|                    | Male   | 192.0   | 37.208333 | 12.049563 | 17.0 | 28.0 | 35.0 |   |
| Asian-Pac-Islander | Female | 346.0   | 35.089595 | 12.300845 | 17.0 | 25.0 | 33.0 |   |
|                    | Male   | 693.0   | 39.073593 | 12.883944 | 18.0 | 29.0 | 37.0 |   |
| Black              | Female | 1555.0  | 37.854019 | 12.637197 | 17.0 | 28.0 | 37.0 |   |
|                    | Male   | 1569.0  | 37.682600 | 12.882612 | 17.0 | 27.0 | 36.0 |   |
| Other              | Female | 109.0   | 31.678899 | 11.631599 | 17.0 | 23.0 | 29.0 |   |
|                    | Male   | 162.0   | 34.654321 | 11.355531 | 17.0 | 26.0 | 32.0 |   |
| White              | Female | 8642.0  | 36.811618 | 14.329093 | 17.0 | 25.0 | 35.0 |   |
|                    | Male   | 19174.0 | 39.652498 | 13.436029 | 17.0 | 29.0 | 38.0 |   |

|                    |        | 75%   | max  |
|--------------------|--------|-------|------|
| race               | sex    |       |      |
| Amer-Indian-Eskimo | Female | 46.00 | 80.0 |
|                    | Male   | 45.00 | 82.0 |
| Asian-Pac-Islander | Female | 43.75 | 75.0 |
|                    | Male   | 46.00 | 90.0 |
| Black              | Female | 46.00 | 90.0 |
|                    | Male   | 46.00 | 90.0 |
| Other              | Female | 39.00 | 74.0 |
|                    | Male   | 42.00 | 77.0 |
| White              | Female | 46.00 | 90.0 |
|                    | Male   | 49.00 | 90.0 |

```
[15]: max_age = grp_r_s.get_group(('Amer-Indian-Eskimo', 'Male'))['age'].max()
print(f"The maximum age of men of Amer-Indian-Eskimo race is {round(max_age)}.")
```

The maximum age of men of Amer-Indian-Eskimo race is 82.

8. Among whom is the proportion of those who earn a lot (>50K) greater: married or single men (marital-status feature)? Consider as married those who have a marital-status starting with Married (Married-civ-spouse, Married-spouse-absent or Married-AF-spouse), the rest are considered bachelors.

```
[16]: # Married men

filt = (df['sex'] == 'Male') & (df['marital-status'].str.startswith('Married'))
df[filt]['salary'].value_counts(normalize = True)
```

```
[16]: salary
<=50K    0.559486
>50K     0.440514
Name: proportion, dtype: float64
```

```
[17]: # Single men

filt = (df['sex'] == 'Male') & ~(df['marital-status'].str.startswith('Married'))
df[filt]['salary'].value_counts(normalize = True)
```

```
[17]: salary
<=50K    0.915505
>50K     0.084495
Name: proportion, dtype: float64
```

9. What is the maximum number of hours a person works per week (hours-per-week feature)? How many people work such a number of hours and what is the percentage of those who earn a lot among them?

```
[18]: max_load = df['hours-per-week'].max()
#df['hours-per-week'].value_counts()
```

```
print(f"Maximum number of hours a person works per week is {max_load}.")
```

Maximum number of hours a person works per week is 99.

```
[19]: num_workaholics = df[df['hours-per-week'] == filt].shape[0]
      print(f"Number of people that work such a number of hours is {num_workaholics}.
      ↪")
```

Number of people that work such a number of hours is 4.

```
[20]: filter1 = (df['hours-per-week'] == df['hours-per-week'].max())
      who_work_for_max_hperw = df.loc[filter1].shape[0]
```

```
[21]: filter2 = (df['hours-per-week'] == df['hours-per-week'].max()) & (df['salary']_
      ↪== '>50K')
      who_work_for_max_hperw_with_hs = df.loc[filter2].shape[0]
```

```
[22]: calc = who_work_for_max_hperw_with_hs / who_work_for_max_hperw * 100
      print(f"Percentage of people who work for maximum no.of weeks with higher_
      ↪salary is {round(calc)}%.")
```

Percentage of people who work for maximum no.of weeks with higher salary is 29%.

10. Count the average time of work (hours-per-week) for those who earn a little and a lot (salary) for each country (native-country). What will these be for Japan?

```
[23]: grp_n_s = df.groupby(['native-country', 'salary'])['hours-per-week'].mean()
      grp_n_s
```

```
[23]: native-country  salary
?                  <=50K    40.164760
                  >50K    45.547945
Cambodia           <=50K    41.416667
                  >50K    40.000000
Canada             <=50K    37.914634
                  ...
United-States      >50K    45.505369
Vietnam            <=50K    37.193548
                  >50K    39.200000
Yugoslavia         <=50K    41.600000
                  >50K    49.500000
Name: hours-per-week, Length: 82, dtype: float64
```

```
[24]: grp_n_s['Japan']
```

```
[24]: salary
<=50K    41.000000
>50K     47.958333
Name: hours-per-week, dtype: float64
```

```
[25]: print(f"Average time of work (hours-per-week) for those who earn a little in_
      ↪Japan is {grp_n_s['Japan'][0]}.")
      print(f"Average time of work (hours-per-week) for those who earn a lot in Japan_
      ↪is {grp_n_s['Japan'][1]}.")
```

Average time of work (hours-per-week) for those who earn a little in Japan is 41.0.

Average time of work (hours-per-week) for those who earn a lot in Japan is 47.958333333333336.


1. How does the salary distribution differ between different marital statuses?

```
[26]: pivot_table_df = df.pivot_table(index='marital-status', columns=['sex',
      ↪'salary'], aggfunc='size', fill_value=0)
      print(pivot_table_df)
```

| sex                   | Female |      | Male  |      |
|-----------------------|--------|------|-------|------|
| salary                | <=50K  | >50K | <=50K | >50K |
| marital-status        |        |      |       |      |
| Divorced              | 2493   | 179  | 1487  | 284  |
| Married-AF-spouse     | 8      | 6    | 5     | 4    |
| Married-civ-spouse    | 903    | 754  | 7381  | 5938 |
| Married-spouse-absent | 194    | 11   | 190   | 23   |
| Never-married         | 4601   | 166  | 5591  | 325  |
| Separated             | 614    | 17   | 345   | 49   |
| Widowed               | 779    | 46   | 129   | 39   |

```
[33]: plt.figure(figsize=(11, 5))
      sns.heatmap(pivot_table_df, cmap="GnBu", annot = True, fmt = ".0f")
```

```
[33]: <Axes: xlabel='sex-salary', ylabel='marital-status'>
```



Heatmap of Marital Status vs Sex and Salary: This heatmap shows the count of records for each


combination of marital-status, sex, and salary. The color intensity in the heatmap represents the count of records. Darker colors represent higher counts.

```
[ ]:
```

2. Create two kdeplots of the hours-per-week feature for each gender on the same chart.

```
[36]: plt.figure(figsize=(10, 5))
sns.kdeplot(data = df, x="hours-per-week", hue = 'sex', multiple = 'stack',
            alpha=.7, linewidth = 0.5, palette = 'crest')
```

```
[36]: <Axes: xlabel='hours-per-week', ylabel='Density'>
```



KDE Plot of Hours per Week by Sex: This plot shows the distribution of hours-per-week for each sex. The KDE (Kernel Density Estimate) plot is a way to estimate the probability density function of a continuous random variable. It uses a Gaussian kernel to estimate the PDF. In this plot, the sex is differentiated by color.

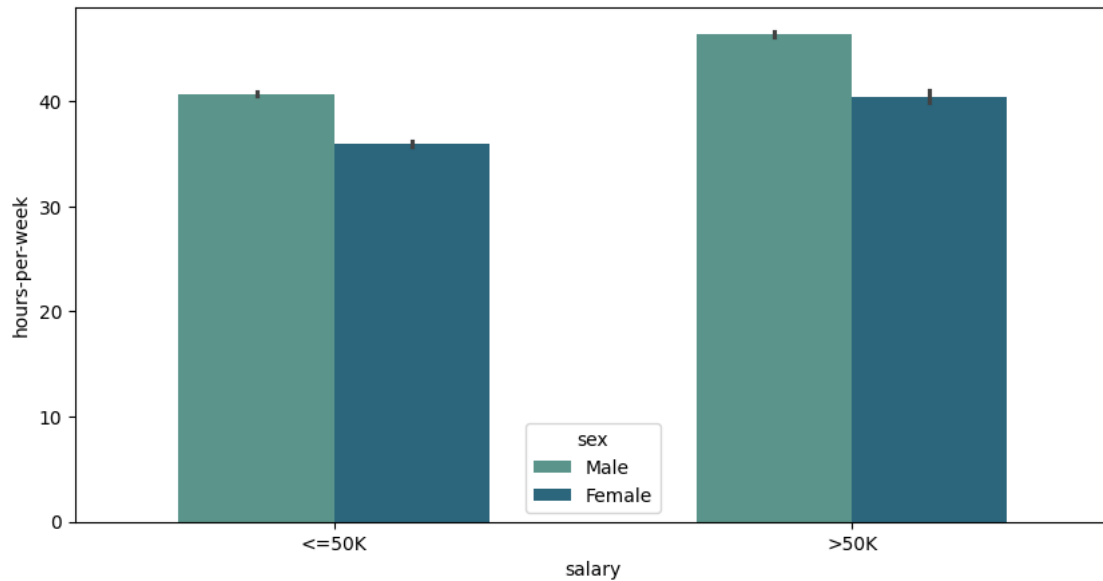
```
[ ]:
```

3. How does the average hours worked per week vary between genders for each salary level?

```
[37]: plt.figure(figsize=(10,5))
sns.barplot(data = df, x = 'salary', y = 'hours-per-week', hue = "sex", width =
            0.6, palette = "crest")
```

```
[37]: <Axes: xlabel='salary', ylabel='hours-per-week'>
```





Bar Plot of Salary vs Hours per Week by Sex: This bar plot shows the average hours-per-week for each salary level, differentiated by sex. The height of the bar represents the average hours-per-week.

[ ]:


[30]: 4. How does the count of records vary with different combinations of age and hours worked per week?

Object `week` not found.

```
[38]: # Define bins for 'age' and 'hours-per-week'
plt.figure(figsize=(11,5))
bins_age = pd.cut(df['age'], bins=range(0, 100, 10))
# bins_age_as_strings = bins_age.astype(str).str.replace(r'[\(\)]', ' ', regex=True).replace(r'[\,\,]', ' - ', regex = True)
bins_hours_per_week = pd.cut(df['hours-per-week'], bins=range(0, 100, 10))
heatmap_data = pd.crosstab(bins_age, bins_hours_per_week)

sns.heatmap(heatmap_data, cmap='crest', annot = True, fmt=".0f")
```

[38]: <Axes: xlabel='hours-per-week', ylabel='age'>




Heatmap of Age vs Hours per Week: This heatmap shows the count of records for each combination of age and hours-per-week bins. The age and hours-per-week are divided into bins of size 10. The color intensity in the heatmap represents the count of records. Darker colors represent higher counts.

[ ]:

5. How does the proportion of salary levels vary within each education group?

```
[40]: fig, ax = plt.subplots(figsize=(11, 5))
df.groupby(['education'])['salary'].value_counts(normalize=True).unstack().
    plot(kind='bar', cmap = "crest", ax=ax)
plt.show()
```



Bar Plot of Education vs Salary: This bar plot shows the proportion of each salary level within each education group. The height of the bar represents the proportion of salary.